

Istio – An introduction for developers

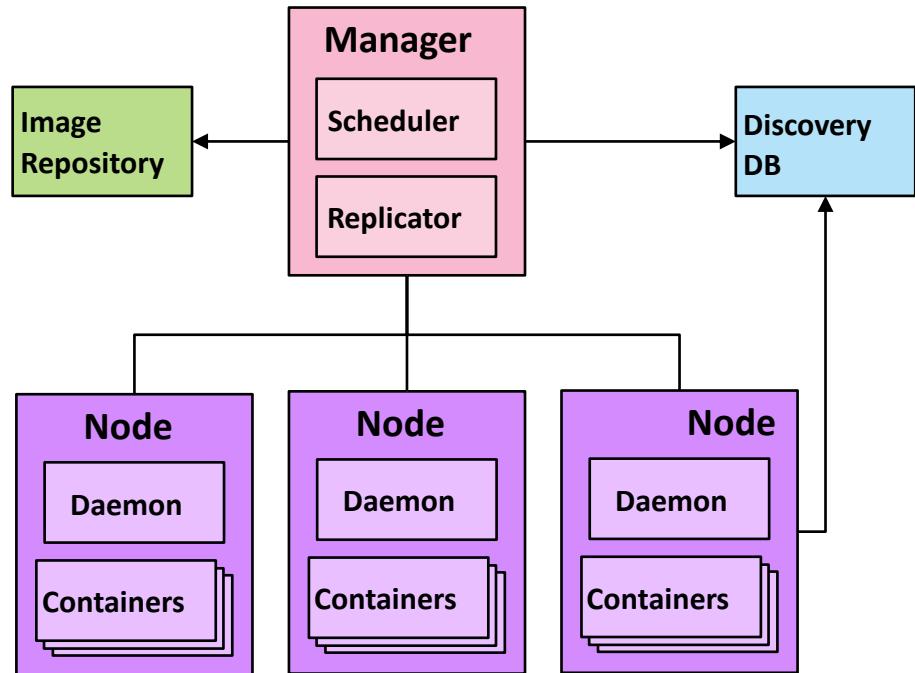
WW Developer Advocacy Team

Benefits of Kubernetes

- Automated scheduling and scaling
- Zero downtime deployments
- High availability and fault tolerance
- A/B Testing



kubernetes



Microservice Challenges

Controlling Traffic

1. How do I do canary testing?
2. How do I A/B testing?

← All of these problems are common across services no matter the runtime.

Resilient Services

1. How do I implement circuit breakers?
2. How do I test faults in the system?
3. How do I limit set rate limits for each service?

← It would be cool if we could solve these problems in a standard way **without changing application code**

Telemetry

1. How can I trace a request through my system of multiple services
2. How can I perform monitoring in a distributed deployment?

Security

1. How can I apply policies across services?

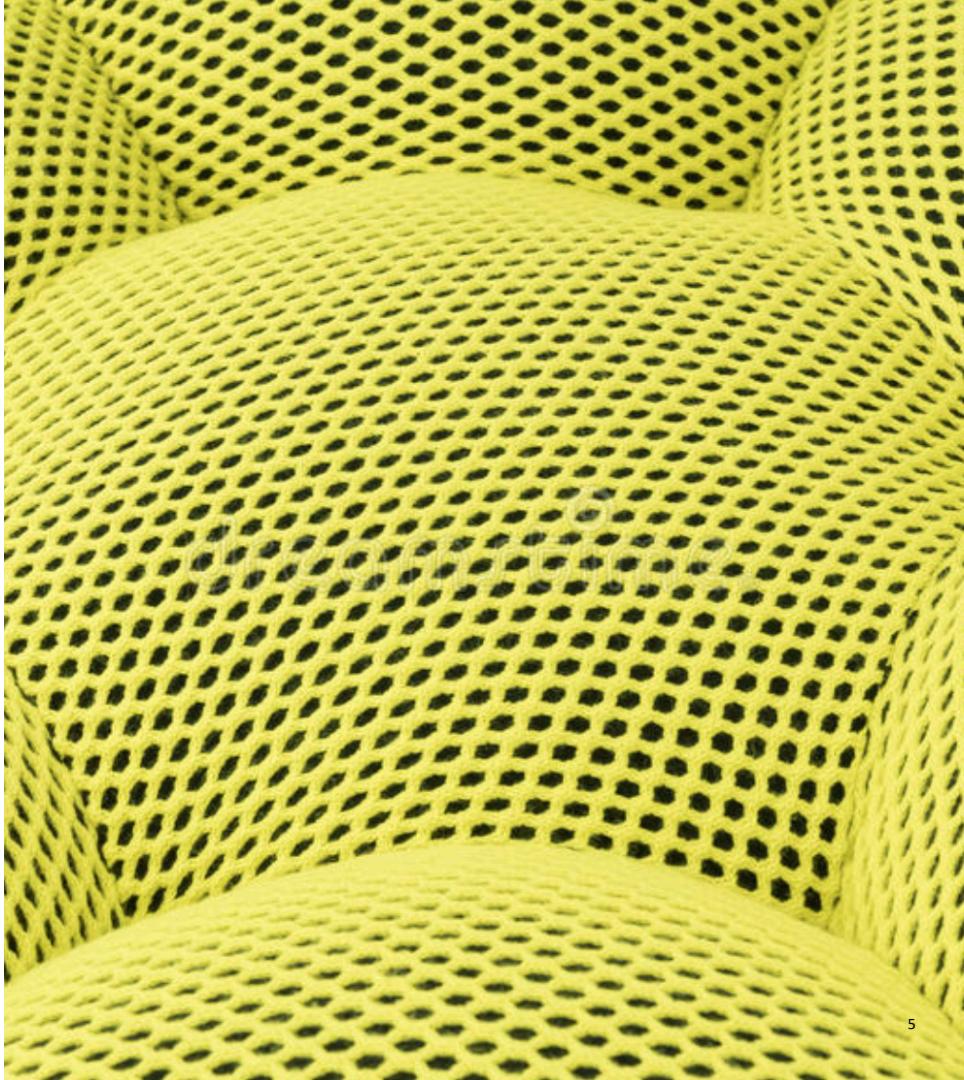
Service Mesh

&

Istio Architecture

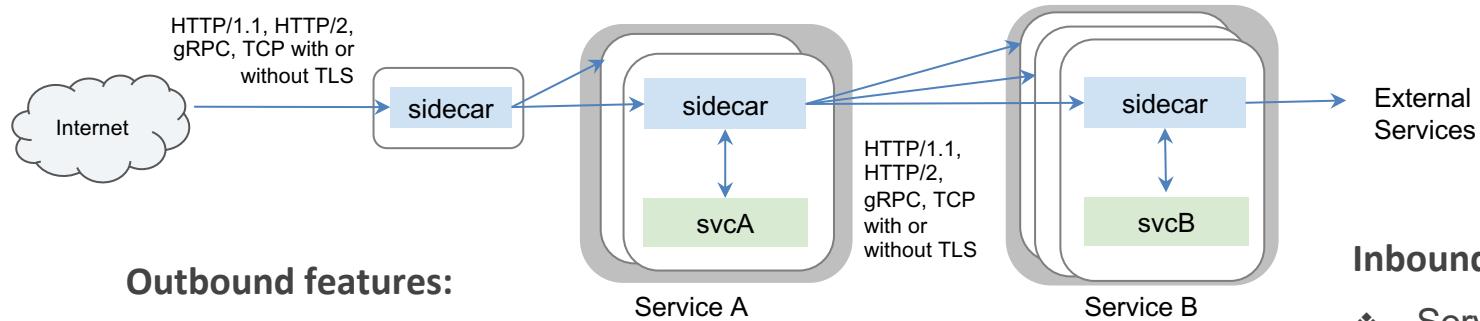
What is a service mesh?

A service mesh provides a **transparent** and **language-independent** way to flexibly and easily manage the **communication** between microservices.



Istio is a **service mesh** that supports managing **traffic** flows between microservices, enforcing **access policies**, and aggregating **telemetry** data, all without requiring changes to the microservice code.

Weaving the mesh



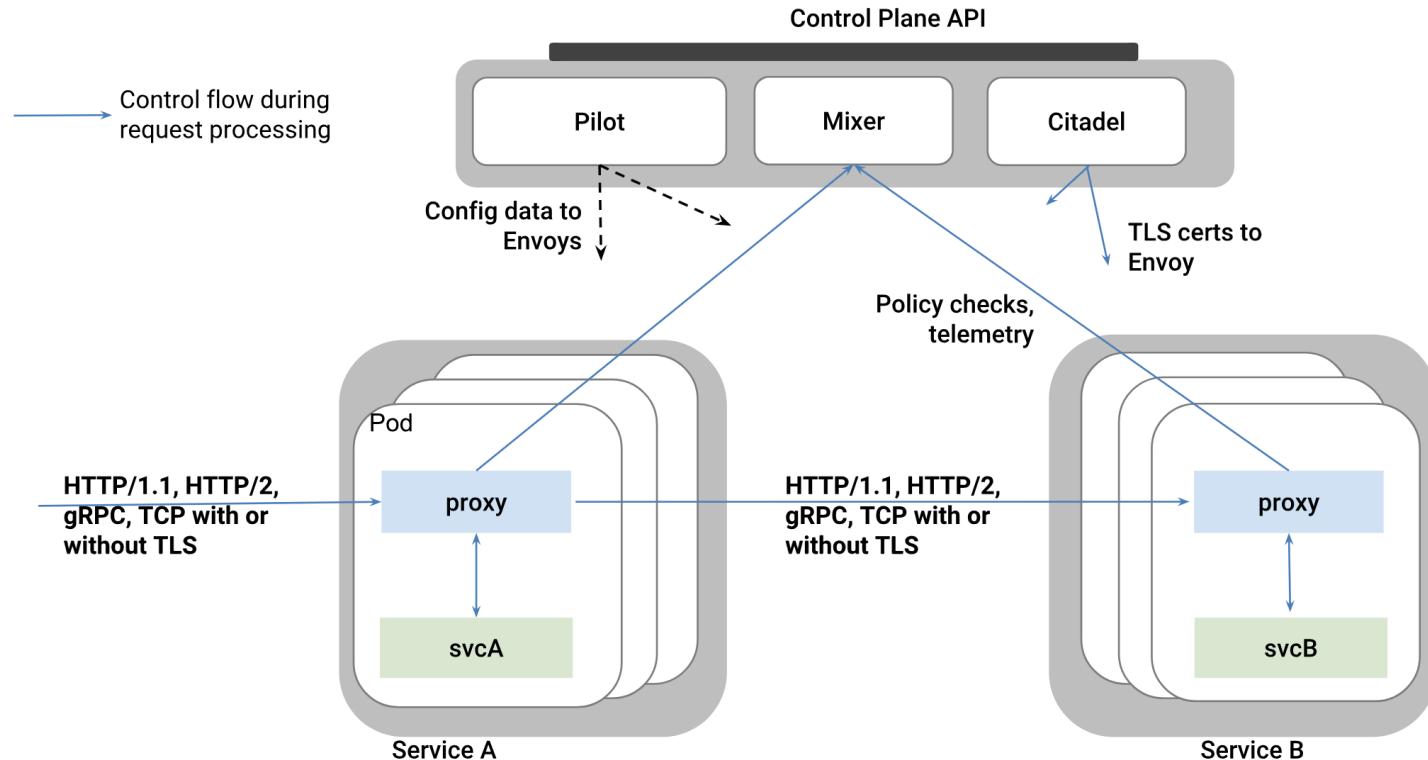
Outbound features:

- ❖ Service authentication
- ❖ Load balancing
- ❖ Retry and circuit breaker
- ❖ Fine-grained routing
- ❖ Telemetry
- ❖ Request Tracing
- ❖ Fault Injection

Inbound features:

- ❖ Service authentication
- ❖ Authorization
- ❖ Rate limits
- ❖ Load shedding
- ❖ Telemetry
- ❖ Request Tracing
- ❖ Fault Injection

Istio Architecture



Istio Architecture

Mixer

- Collects telemetry data sent to it from the proxy sidecar containers on each request.
- Platform-independent, very pluggable with various backends.

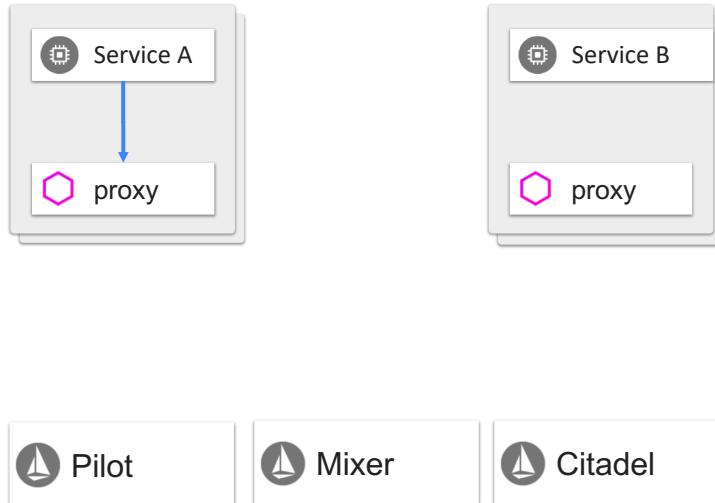
Pilot

- “Drives” the service mesh.
- Listens for your configuration around traffic routing, circuit-breakers and fault injection.
- Converts that to low-level routing rules, and distributes those to sidecar containers at runtime.

Citadel

- Does security
- Key and certificate management. Distributes as Kubernetes Secrets.

Life of a request

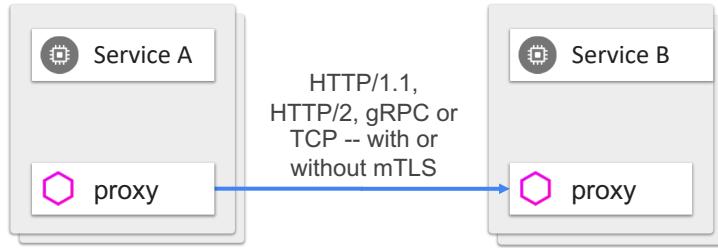


Service A places a call to <http://service-b>.

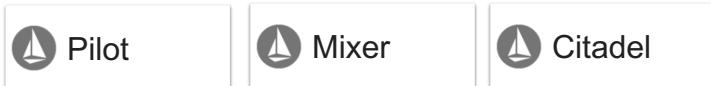
Client-side Envoy intercepts the call.

Envoy consults config (previously received from Pilot) to know how/where to route call to service B (taking into account service discovery, load balancing, and routing rules), forwards the call to the right destination.

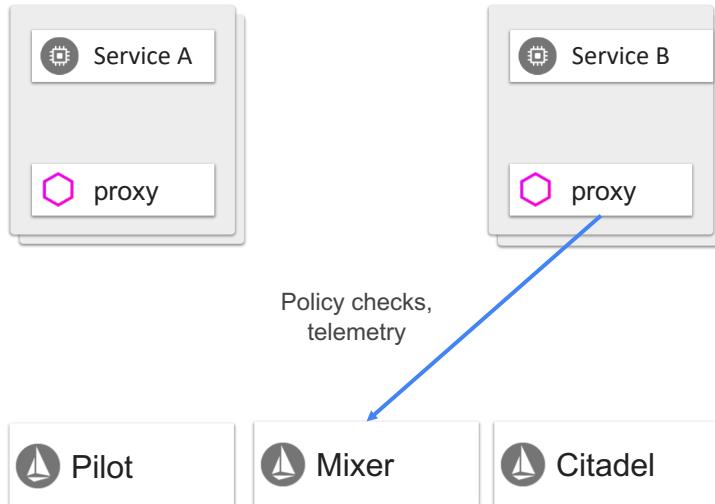
Life of a request



Envoy forwards request to appropriate instance of service B. There, the Envoy proxy deployed with the service intercepts the call.

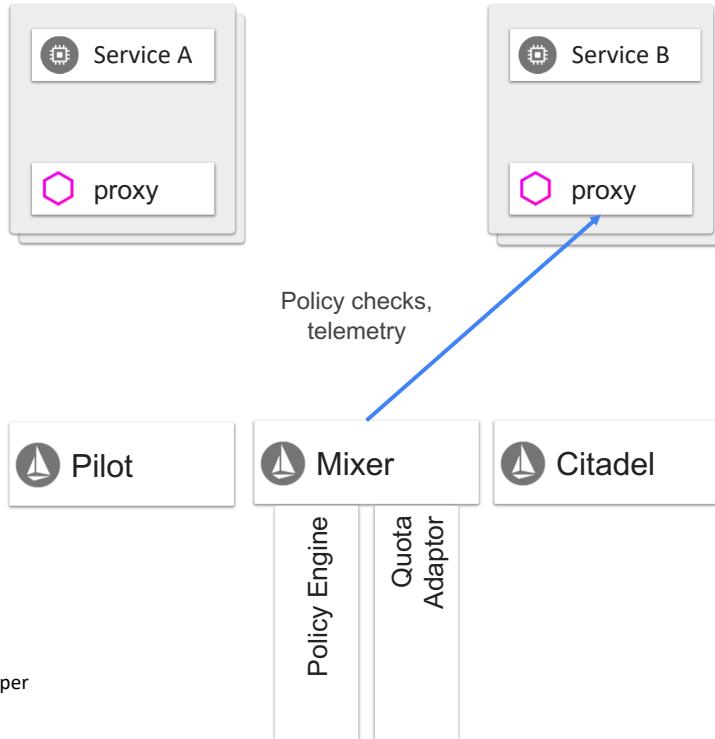


Life of a request



Server-side Envoy checks with Mixer to validate that call should be allowed (ACL check, quota check, etc).

Life of a request

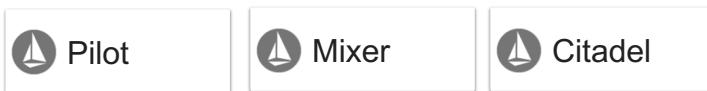


Mixer checks with appropriate adaptors (policy engine, quota adaptor) to verify that the call can proceed and returns true/false to Envoy

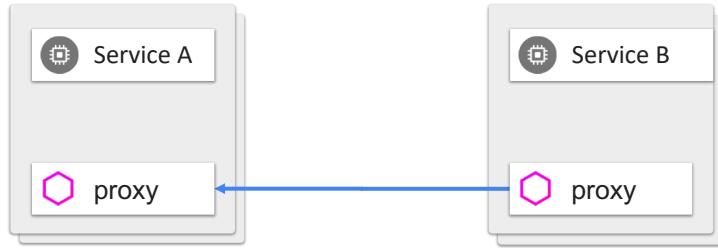
Life of a request



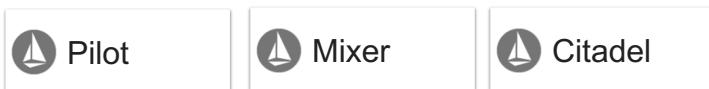
Server-side Envoy forwards request to service B, which processes request and returns response



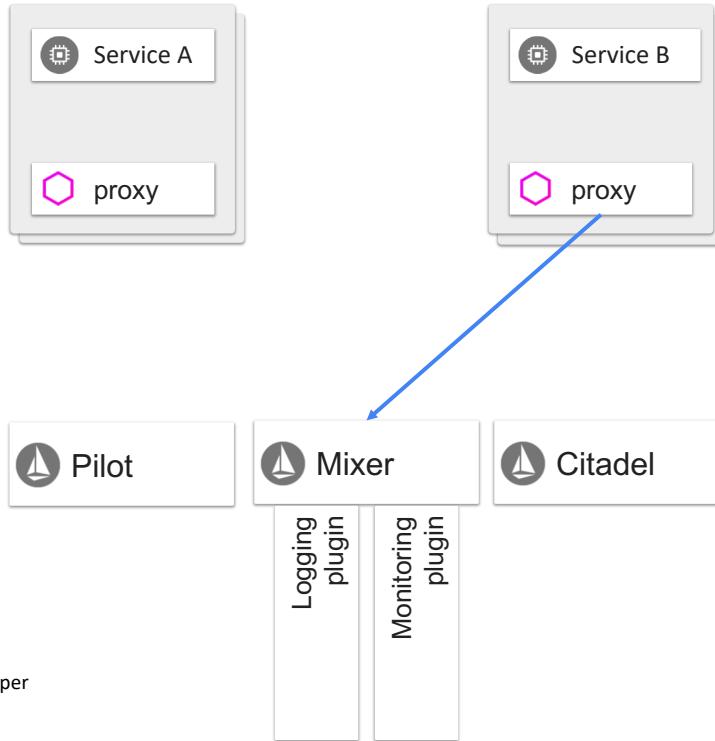
Life of a request



Envoy forwards response to the original caller, where response is intercepted by Envoy on the caller side.



Life of a request

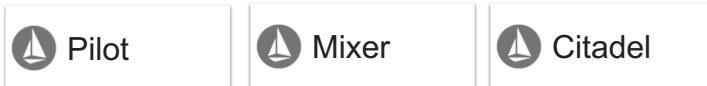


Envoy reports telemetry to Mixer, which in turn notifies appropriate plugins

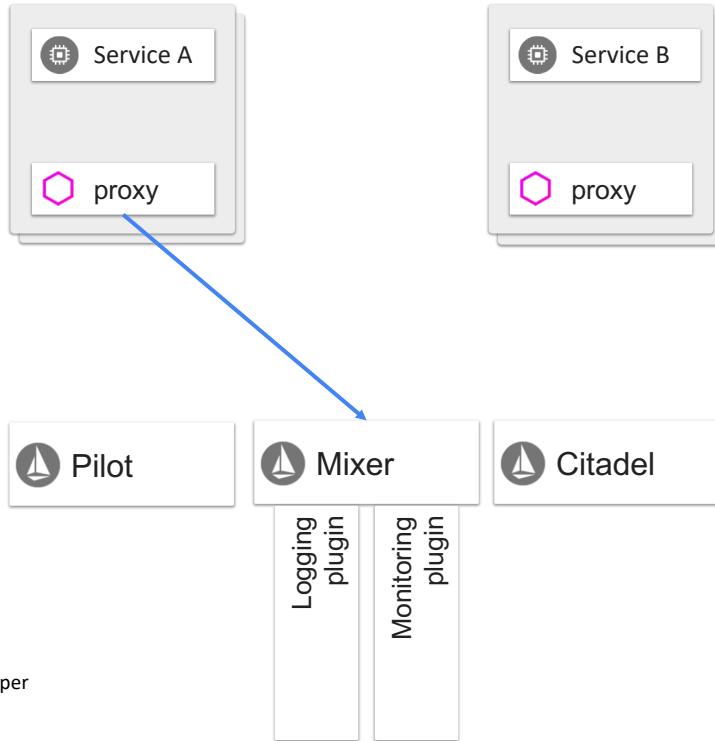
Life of a request



Client-side Envoy forwards response to original caller.

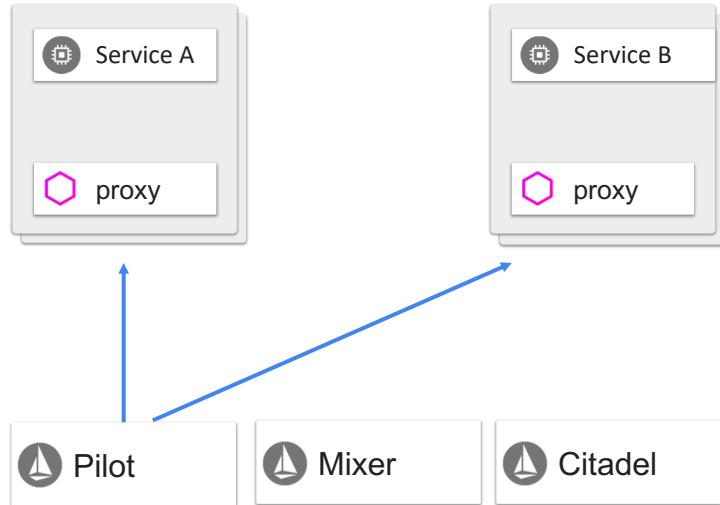


Life of a request



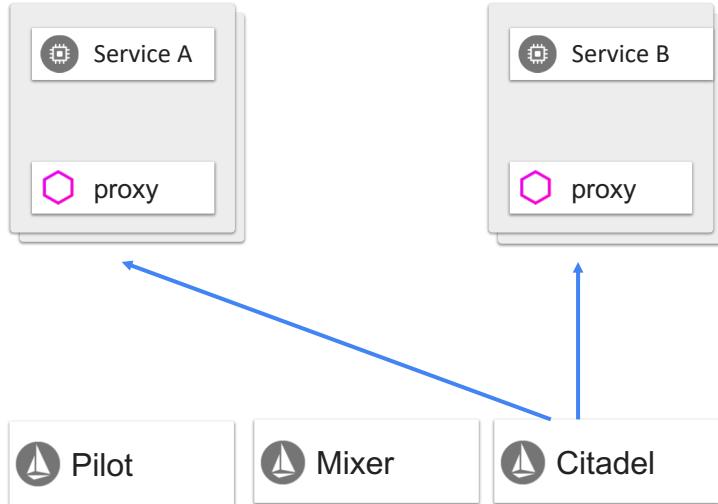
Client-side Envoy reports telemetry to Mixer (including client-perceived latency), which in turn notifies appropriate plugins

At Anytime



Pilot listens to your configuration, such as routing rules, circuit breaking and fault injection. Converts to low-level config (routing rules) and distributes to proxy side cars

At Anytime

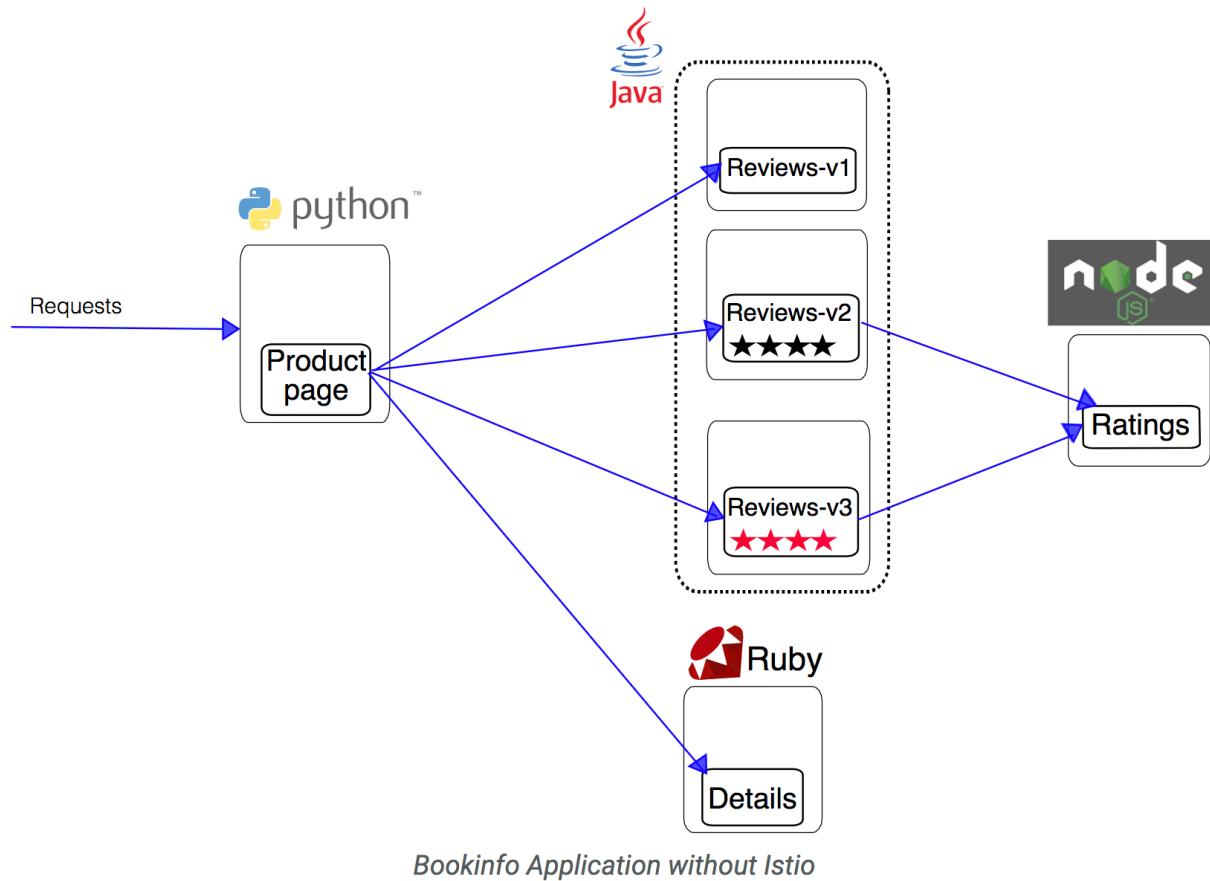


Citadel distributes keys and certificates as Kubernetes Secrets available to sidecar containers

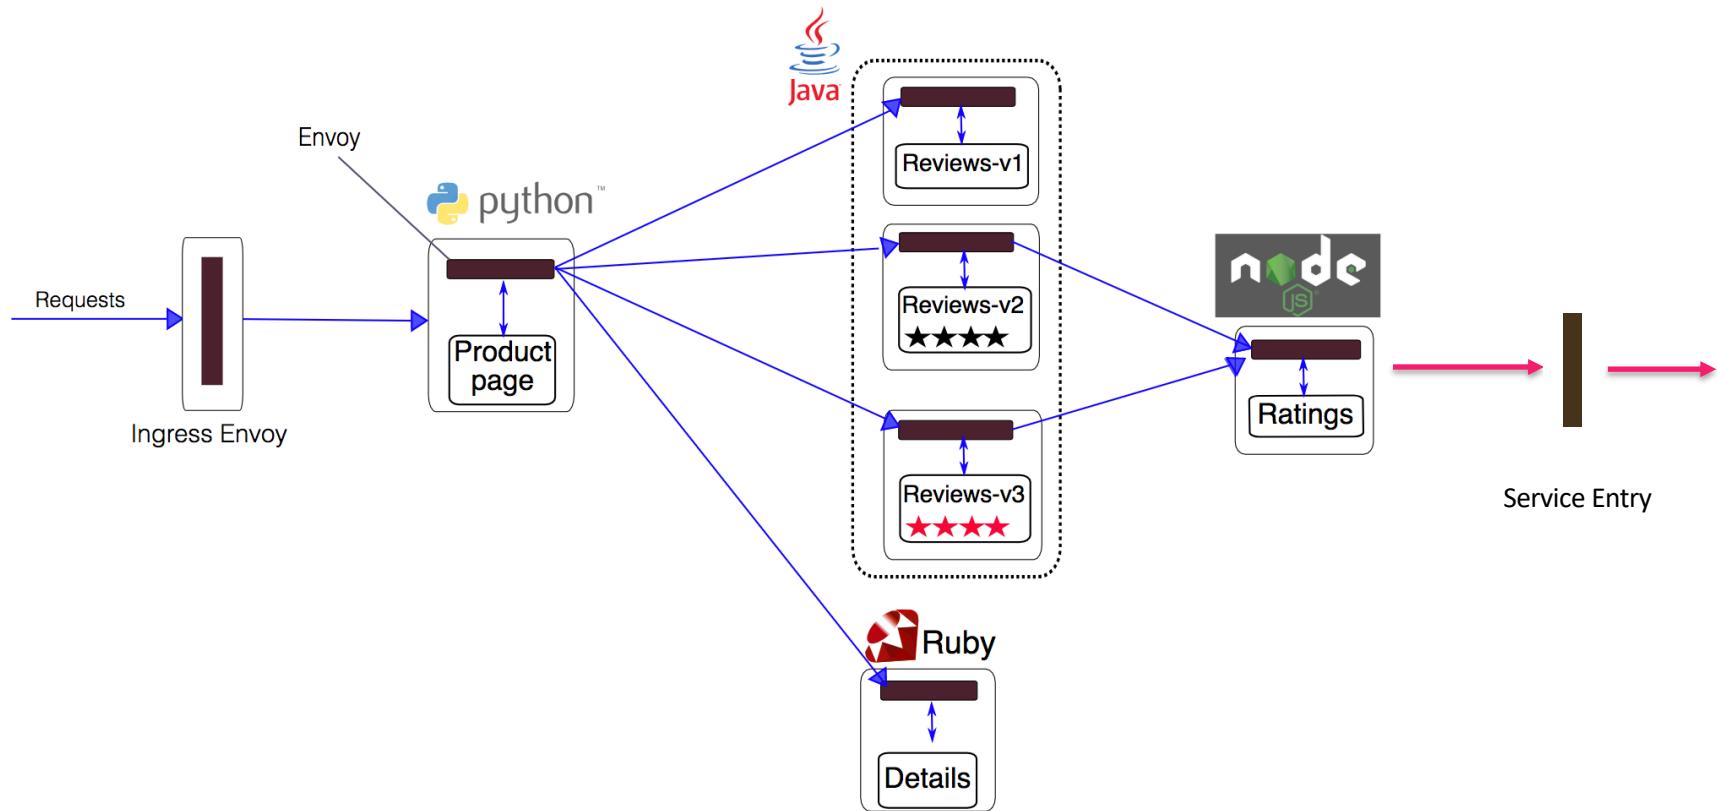
Istio Proxy Sidecar

- To create a service mesh with Istio, you update the deployment of the pods to add the Istio Proxy (based on the Lyft Envoy Proxy) as a side car to each pod. The Proxy is then run as a separate container that manages all communication with that pod.
 - Manual Approach
 - **Automated Approach**
- Istio modifies the individual pods and not deployments

Bookinfo Sample Microservices Application (without Istio)



Bookinfo Sample Microservices Application (with Istio)



Control Ingress Traffic

- **Gateway** describes a load balancer operating at the edge of the mesh receiving incoming or outgoing HTTP/TCP connections.
- **VirtualService** defines a set of traffic routing rules to apply when a host is addressed. Each routing rule defines matching criteria for traffic of a specific protocol.
- Istio ingress controller gets exposed as a normal kubernetes service load balancer on port 80.

```
kubectl get svc -n istio-system
```

<https://istio.io/docs/tasks/traffic-management/ingress/>

Istio Examples

Request Routing

```
1  apiVersion: networking.istio.io/v1alpha3
2  kind: VirtualService
3  metadata:
4    name: reviews
5  spec:
6    hosts:
7      - reviews
8    http:
9      - match:
10        - headers:
11          - end-user:
12            exact: jason
13        route:
14          - destination:
15            host: reviews
16            subset: v2
17      - route:
18        - destination:
19          host: reviews
20          subset: v1
```

Canary Testing

Route user:jason to reviews:v2
Others still get reviews:v1

Traffic Shifting

```
1  apiVersion: networking.istio.io/v1alpha3
2  kind: VirtualService
3  metadata:
4    name: reviews
5  spec:
6    hosts:
7      - reviews
8    http:
9      - route:
10        - destination:
11          host: reviews
12          subset: v1
13          weight: 50
14        - destination:
15          host: reviews
16          subset: v3
17          weight: 50
```

50% -> v1

50% -> v3

Rate Limits

```
1  apiVersion: "config.istio.io/v1alpha2"
2  kind: memquota
3  metadata:
4    name: handler
5    namespace: istio-system
6  spec:
7    quotas:
8      - name: requestcount.quota.istio-system
9        maxAmount: 5000
10       validDuration: 1s
11       # The first matching override is applied.
12       # A requestcount instance is checked against override dimensions.
13       overrides:
14         # The following override applies to 'ratings' when
15         # the source is 'reviews'.
16         - dimensions:
17           destination: ratings
18           source: reviews
19           maxAmount: 1
20           validDuration: 1s
21           # The following override applies to 'ratings' regardless
22           # of the source.
23           - dimensions:
24             destination: rating
25             maxAmount: 100
26             validDuration: 1s
```

5000 requests per 1s
ratings: 100 requests per 1s

Circuit Breaking

```
1  apiVersion: networking.istio.io/v1alpha3
2  kind: DestinationRule
3  metadata:
4    name: httpbin
5  spec:
6    host: httpbin
7    trafficPolicy:
8      connectionPool:
9        tcp:
10       maxConnections: 1
11
12      http:
13        http1MaxPendingRequests: 1
14        maxRequestsPerConnection: 1
15      outlierDetection:
16        consecutiveErrors: 1
17        interval: 1s
18        baseEjectionTime: 3m
19        maxEjectionPercent: 100
```

Max 1 concurrent
connection & request

Delay Injection

```
1  apiVersion: networking.istio.io/v1alpha3
2  kind: VirtualService
3  metadata:
4    name: ratings
5  spec:
6    hosts:
7      - ratings
8    http:
9      - match:
10        - headers:
11          |   end-user:
12          |     exact: jason
13        fault:
14          delay:
15            percent: 100
16            fixedDelay: 7s
17        route:
18          - destination:
19            host: ratings
20            subset: v1
21      - route:
22        - destination:
23          host: ratings
24          subset: v1
```

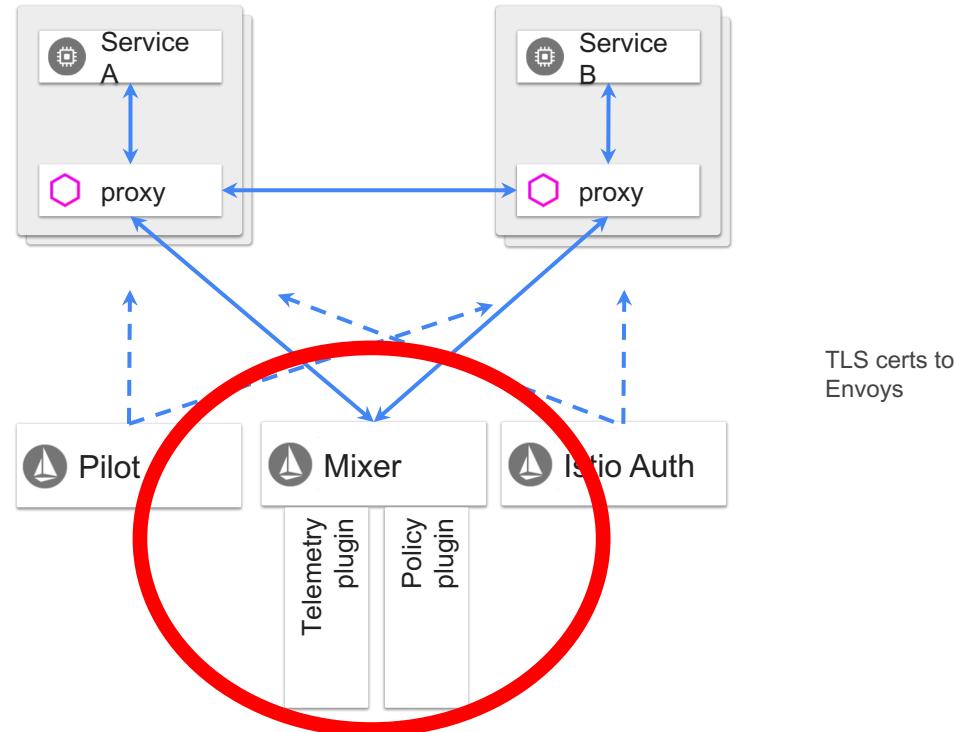
Inject 7 second delay

Fault Injection

```
1  apiVersion: networking.istio.io/v1alpha3
2  kind: VirtualService
3  metadata:
4    name: ratings
5  spec:
6    hosts:
7      - ratings
8    http:
9      - match:
10        - headers:
11          |   end-user:
12          |     exact: jason
13        fault:
14          abort:
15            percent: 100
16            httpStatus: 500
17        route:
18          - destination:
19            host: ratings
20            subset: v1
21      - route:
22        - destination:
23          host: ratings
24          subset: v1
```

jason: Return with Error 500

Istio Telemetry (v1)



Telemetry Tools

Metrics



Service Mesh Visualization



Dashboard for Metrics



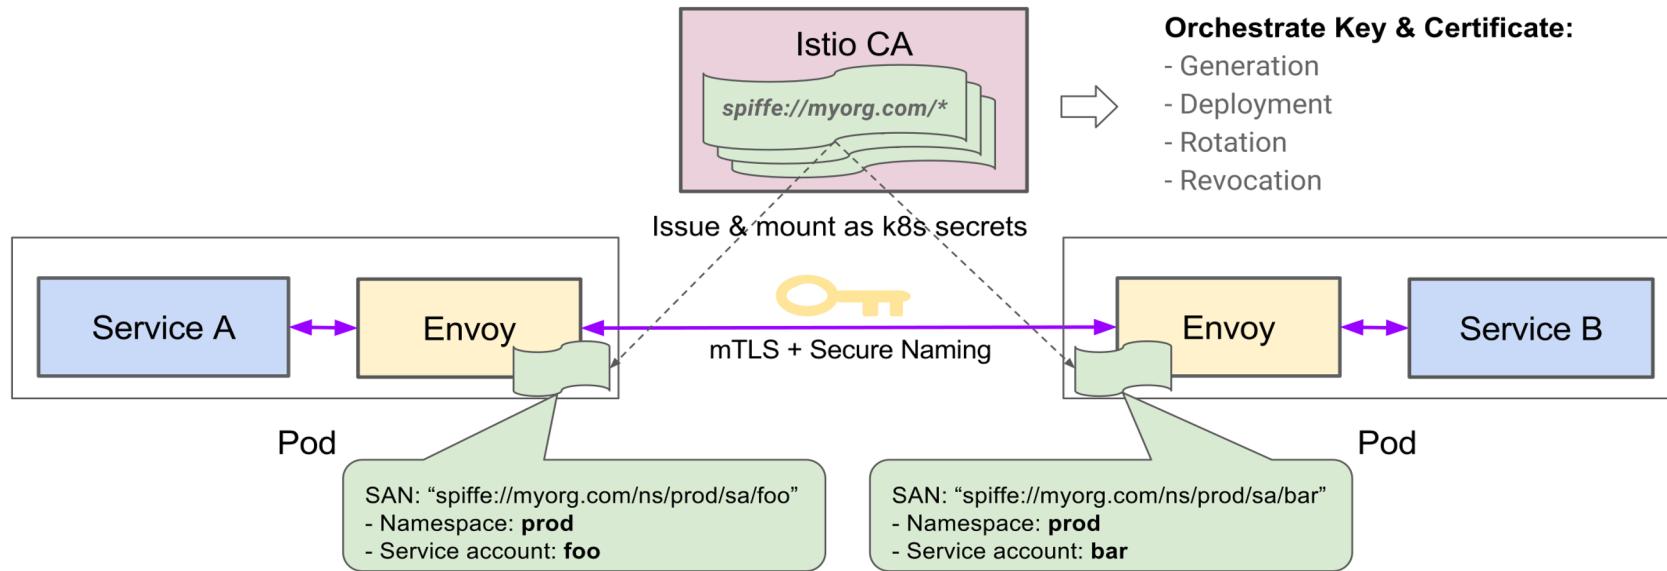
Request Tracing



Securing Microservices

- Verifiable identity
- Secure naming / addressing
- Traffic encryption
- Revocation

Istio Citadel



Resources

Istio documentation

<https://istio.io/docs/>

Istio in IBM products

<https://www.ibm.com/cloud/info/istio>

The screenshot shows the IBM Cloud interface. At the top, there's a navigation bar with 'IBM Cloud' and a search bar. Below the navigation bar, the user is in the 'Clusters' section, specifically viewing the 'user009-cluster'. On the left, there's a sidebar with 'Access', 'Overview', 'Worker Nodes', 'Worker Pools', and 'Add-ons'. The 'Add-ons' tab is currently selected and highlighted with a blue border. Under 'Add-ons', there are two items: 'Managed Istio v1.4' and 'Managed Knative Experimental'. The 'Managed Istio' card is expanded, showing its status as 'Normal - Addon Ready' and a brief description: 'Istio is a service mesh for providing a uniform way to integrate microservices, manage traffic flow across microservices, enforce policies and aggregate telemetry data. The Istio add-on simplifies the installation and maintenance of your istio control plane so you can focus on managing your microservices.' A 'Learn More' link is also present. To the right of the 'Managed Istio' card is a 'Kubernetes dashboard' button and an 'Actions...' dropdown menu. The 'Managed Knative' card is partially visible below it, with an 'Install' button.

Lab

<https://ibm.biz/BdzP4i>

